
Physical Attacks and Countermeasures

vincent.immler@oregonstate.edu

CS579/ECE599, January 12, 2023

SISD MISD fault tolerance
SIMD MIMD multiple system



Oregon State
University

Introduction to Cryptography (Engineer's Perspective)

- Note: this is only intended as brief review of concepts that matter
- Classical scenario: Alice and Bob are honest parties
- The communication link between them is not secure
- They want to communicate *securely*
 - By using the same secret key: symmetric cryptography
 - By using public/private keys: asymmetric cryptography
- Eve can eavesdrop the communication
- Most security principles have been formulated with this in mind

Security Principles of Cryptography

- Kerckhoff's principle (1883): *"the security of a cryptosystem must lie in the choice of its keys only; everything else (including the algorithm itself) should be considered public knowledge."*
 - It is easier to keep/change a secret key than a secret algorithm
 - Asymmetric schemes and communication with many parties are possible
 - Public scrutiny allows to inspect security and avoid backdoors
- *"A cryptographic scheme for a given task is secure, if no adversary of a specified power can achieve a specified break"* [Katz and Lindell]
- Issues of above definitions within the context of physical security
 - 'Cryptosystem' or 'scheme' clearly have an algorithm or protocol in mind
 - What about cryptographic devices that implement these?
 - What about non-cryptographic security objectives?
- The latter is more in line with nowadays security certifications (Common Criteria)
- Breaking a chip or cryptographic device means violating its security objective(s)

Security Principles of Modern Cryptography

- Cryptographic security principle: computational complexity
 - Most widely used; complexity expressed in terms of bits
 - Exhaustive key search (brute force) of key is not feasible
 - Symmetric ciphers typically require ≥ 80 bit for security
 - Asymmetric ciphers rely on mathematical trapdoor function
 - Assumption: inverting the trapdoor is difficult (no good algorithms known)
 - Security level depends on mathematical problem; risk of quantum computers
- Cryptographic security principle: indistinguishability
 - Attacker is given output of valid cryptographic protocol and a random output
 - Attacker is tasked to distinguish between the two
 - Attacker's advantage should be negligible
 - Security proof of various cryptographic construction rely on this 'game'
- The concept of indistinguishability is important for physical security, too

Security Principles of Information Theory

- Fundamental topics of information theory include coding theory
 - Error Correcting Codes (ECC) and error detection
 - Source coding and compression
 - Post-Quantum Cryptography (PQC) often relies on coding theory
- Information-theoretic security principles
 - Communication link is often modeled as “Wyner’s Wiretap Channel”
 - To eavesdrop on channel, Eve is limited by *noise* → *physical* assumption!
- Please note: bits of security are not made all equal:
 - In computational complexity: Shannon entropy bits (‘on average’)
 - In information-theoretic context: Shannon, or *min*-Entropy, or ...
- Relevant for Physical Unclonable Functions (PUFs), Side-Channel Analysis (SCA)
- Comparison of cryptographic and information-theoretic notions here ¹

¹“Semantic Security for the Wiretap Channel” by Bellare et al. (2012)

Concept of Trust vs. Security

- Trust: a complicated notion related to belief in honesty, reliability, etc. of an entity
- Trusted: a system or component whose failure can break the security policy
- Trustworthy: the degree to which the behavior of the component is demonstrably compliant with its stated functionality
- Demonstrably compliant = having assurance
- A system is defined as *trustable* as long as “*if it always behaves in the expected manner for the intended purpose*” [Trusted Computing Group]



(source: Google)

Motivation for Physical Attacks

- Hack once, own them all, possibly forever
 - Extract firmware or CSPs from one device; own all the other devices
 - Enable subsequent software/firmware analysis
 - Extend functionality, product piracy, to show off, etc.
- Hack what is outside the software scope
 - Low level attack = difficult to counteract in software
 - Increasingly more functionality in hardware, hidden from software
 - Example: actual key storage and encryption in hardware; only API call
- Lack of feedback to manufacturer
 - Device can be attacked without anyone knowing (e.g., unlike websites)
 - No monitoring etc. outside of what shipped within the device
 - Therefore no need to disclose the applied technique!
- Software security 'just' a logical problem
 - Hardware-assisted schemes strong, e.g., pointer authentication, tagged memory
 - Exploits working around these a million dollar business → proof of difficulty

In General: Physical Attack Scenarios and Goals

- Attacker can be legitimate owner of the device
 - Can use device in intended ways
 - Can manipulate, disturb, and observe the device
- Attack may target CSPs, or firmware, etc.
- Revealed information may be used for:
 - Counterfeit product or to modify it in unintended ways
 - Gain unauthorized access to sensitive information
 - Possibly tap into the communication of others
- Physical attack means to use physical properties of the device to break it
- Therefore, typically requires control over the device or close proximity

Quick Facts: Timing Side-Channel Analysis

- The term “Timing Attack” was first introduced at CRYPTO’96 by Paul Kocher (back then: Cryptography Research Inc., today: a division of Rambus)
- Few other theoretical approaches without practical experiments; first publicly known practical results in early 1998
- Basic observation: processing time may depend on the data processed
- Absolute duration or differences can be measured
- Practical requirements for attack:
 - Possibility to (indirectly) monitor processed secret data
 - Have a way to automate and record the process
 - Basic computational and statistical tools
 - Either know or suspect properties of the implementation

Example Attack: Electronic Safe Lock

- Suppose you have a safe with an electronic lock with PIN entry
 - In order to open the safe, a secret 8-digit PIN must be entered (0...9)
 - Each digit is represented by one byte
 - The device has a delayed response upon verification of 100 ms
 - For a brute-force attack: there are 10^8 possible PIN codes
 - Trying all possible PINs (without entering them) already takes 116 days
- Formally secure enough ...but are there better attacks?



(source: lock I recently bought on ebay)

Example Attack: How to Approach

- Suppose that attacker has access to a reference lock
 - Learns the type of 8-bit microcontroller used and its instruction set
 - Is able to automate the PIN entry
- Let us assume the PIN check is done digit by digit (as shown below)
- Early return on wrong digit leads to data-dependent process
- How to attack? Divide-and-conquer strategy!

```
1 unsigned char verifyPIN(unsigned char Pref[], unsigned char P[]) {  
2     for(int i=0; i<7; i++) {  
3         if (P[i] != Pref[i])  
4             return 0xEE; // return fail code  
5     }  
6     return 0x00; // return pass code  
7 }
```

Example Attack: Timing Side-Channel Analysis

- Due to automation, attacker is able to measure time
 - Time between PIN entry to response matters
 - Response can be audible (buzzer), visible (LED), or wire tapped
 - Response detection can be automated
- Attacker selects a fixed PIN and changes only first digit
 - 9x the wrong result (same timing)
 - 1x the correct PIN-digit (different timing)
 - Difference in occurrence gives away important information
 - Likely the correct PIN-digit has a longer response (return on 2nd digit)
- After finding first digit, remaining digits work the same
- Procedure only takes $10 \times 8 = 80$ timing measurements!

Example: Possible Countermeasures and Obstacles

- Timing cannot be measured precisely → multiple measurements and averaging
- PIN retry counter with additional delay → attack slowed down but still works
- PIN retry counter with hard limit of X attempts
 - Consider the business scenario: people forget their PIN and may try too often
 - Fault attacks may help to either skip counter increment or reset to previous state
- You are an uber-programmer and implement an all-digits-at-once comparison
 - Remember, \$\$\$ is important and your boss only approved an 8-bit microcontroller
 - In this case, lack of an all-digit-at-once comparison has little to do with your skills
- Developer of system is constrained by time, money, marketing, physics, ...
- Attacker possibly limited by laws of physics only!
- Considering the cost of attacks vs. product difficult

Example Attack: Continued

- Requirement: Runtime of PIN verification must be independent of entered PIN
 - Not easy but possible, e.g., balancing execution time by NOP instructions
 - Other microarchitectural problems: caches, out-of-order execution, etc.
 - Of course, an 8-bit microcontroller has a very simple microarchitecture
- Let's assume timing behavior has been fixed – problem solved?
- No – most digital circuits are implemented by CMOS logic
 - Power consumption depends on executed instruction and processed data
 - Part of this course to study this type of attack in greater detail
 - For now, we take this for granted (without understanding the 'why')

Example Attack: Power Analysis

- Instead of measuring time, attacker measures power consumption
- Attacker selects a fixed PIN and changes only first digit
 - 9x the wrong result (same power signature)
 - 1x the correct PIN-digit (different power signature)
 - Difference in occurrence tells you something
 - Likely the correct PIN-digit has most significant difference
- After finding first digit, remaining digits work the same
- Procedure only takes $10 \times 8 = 80$ timing measurements!
- Practical limitations such as measurement noise make this more difficult

Example: Possible Countermeasures and Obstacles

- Solving this dependency of processed data and executed path is difficult
 - Custom logic style needed in hardware (even then, not easy)
 - Masking the actually processed data “somehow”
- Arbitrary timing / random delays make this more difficult, too
- Random processing of digits upon verification (“shuffling”)
- Doing something random requires ...randomness!
- True Random Number Generator (TRNG) often not included
- Switched capacitor design (draw power from internal buffer)
- Consider this all solved – are we still able to defeat the lock?

Example: Additional Attacks

- Other ways to deduce similar side-channel information
 - Electromagnetic Emanation (EM)
 - Thermal emission / Temperature
 - Photonic emission
 - Acoustics
- Attacker may not be concerned about learning the PIN
 - Tampering with frequency of circuit (clock glitch) to skip verification
 - Tampering with voltage supply (power glitch) to skip verification
 - For the attacker, only one solution is enough to bypass security
- Why bother with the PIN? A spiker/zapper may initiate bolt movement
- Alternative approach: extract recovery PIN from a reference device

Implementations Matter!

- By way of example, we have seen that
 - Securely implementing a PIN verification is difficult
 - Improper programming was only one of the ways to fail
 - Not knowing processor related dependencies was yet another
 - Perhaps we cared about the wrong part of the system
 - We were not even limited by project deadlines or business aspects
- Non-standard operations: difficult to solve satisfactorily
- Standard operations such as encryption best implemented in hardware
- We have to learn a little bit of crypto to understand the implementation
- NEVER implement security on your own; always have it tested!

Applied Cryptography

Advanced Encryption Standard

Advanced Encryption Standard (AES)

- Symmetric ‘block’ cipher
- Successor of Data Encryption Standard (DES)
- Selection through open competition
- Selection by US National Institute of Standards and Technology (NIST)
- Likely to be the dominant cipher for the next two decades
- Much more secure and efficient than DES or 3DES
- cf. NIST FIPS 192 “Advanced Encryption Standard” ²

²<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.197.pdf>

AES Timeline

- AES competition announced in January, 1997
- 15 candidates submitted by August, 1998
- 5 finalists announced in August 1999:
 - Mars, IBM Corporation
 - RC6, RSA Laboratories
 - Rijndael, Joan Daemen and Vincent Rijmen
 - Serpent, Eli Bihman et al.
 - Twofish, B. Schneier et al.
- Rijndael standardized as AES on October, 2001

Bruce Schneier Facts

Things you might not know about Bruce Schneier

[Contact](#) [Random RSS Feed](#) [Top 10 Facts](#) [Suggest Fact](#)




[← Previous Fact](#) | [Random Fact](#) | [Latest Fact](#) | [Search Facts](#) | [Next Fact →](#)

AES stands for "Ain't Encryption to Schneier."

(source: schneierfacts.com)

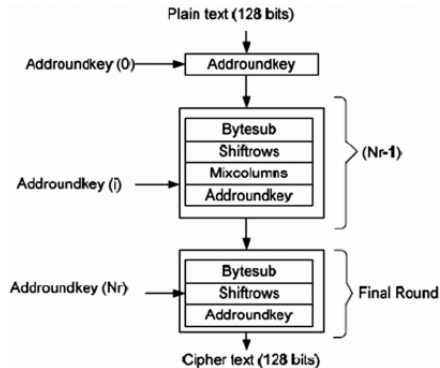
High-level View on AES

- AES supports 3 key lengths: 128, 192, 256 bit
- 128 bit key already resistant against brute-force attack for foreseeable future
- No relevant weaknesses despite 20+ years of research in the crypto community
- NSA allows AES for classified documents up to TOP SECRET; similar to other countries




AES Basic Structure

- AES comprised of a Substitution Permutation Network (SPN) repeated for several rounds
- The block size and data path are 128 bit
- Number of rounds: 10, 12, 14 based on selected key length 128, 192, 256
- Notable components:
 - Key Addition: XOR with roundkey
 - Byte substitution: 8x8 S-Boxes
 - Diffusion Layer: create an avalanche effect by spreading influence of each input bit over many output bits
 - Here: diffusion layer created by ShiftRows and MixColumns



AES AddRoundKey and SubBytes

- 4x4 byte-wise matrix to represent internal state
- All fundamental operations with this matrix
- AddRoundKey is a bitwise XOR operation
 - Applies to full width of data path
- SubBytes operates on each byte independently
 - all Sboxes are the same
 - Sbox is inversion in $GF(2^8)$
 - followed by affine transformation



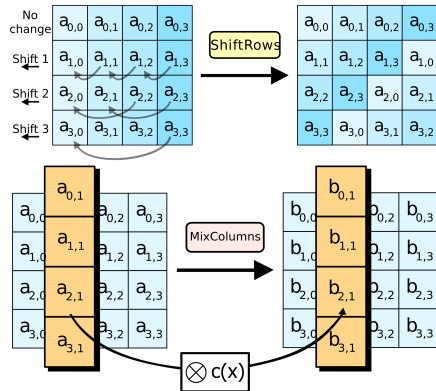
AES ShiftRows and MixColumns

■ ShiftRows

- Rotates each row by 0, 1, 2, 3 bytes to the left
- In Software: adjust index
- In Hardware: just (re-)wiring

■ MixColumns

- Columns represent quarter-state of AES
- Each column is multiplied with fixed matrix
- See next slides for details



Closer Look at MixColumns

- MixColumns is a matrix multiplication as follows:

$$\begin{bmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \end{bmatrix} = \begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \times \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{bmatrix}$$

- Resulting in:

$$b_0 = 02a_0 + 03a_1 + 01a_2 + 01a_3$$

$$b_1 = 01a_0 + 02a_1 + 03a_2 + 01a_3$$

$$b_2 = 01a_0 + 01a_1 + 02a_2 + 03a_3$$

$$b_3 = 03a_0 + 01a_1 + 01a_2 + 02a_3$$

- Since the elements are in $\text{GF}(2^8)$, “special” math applies here
- Plus sign is simply XOR; Multiply by 2 and by 3 needs to be implemented

Multiply by 2 and 3 (from FIPS 197)

4.2.1 Multiplication by x

Multiplying the binary polynomial defined in equation (3.1) with the polynomial x results in

$$b_7x^8 + b_6x^7 + b_5x^6 + b_4x^5 + b_3x^4 + b_2x^3 + b_1x^2 + b_0x. \quad (4.4)$$

The result $x \bullet b(x)$ is obtained by reducing the above result modulo $m(x)$, as defined in equation (4.1). If $b_7 = 0$, the result is already in reduced form. If $b_7 = 1$, the reduction is accomplished by subtracting (i.e., XORing) the polynomial $m(x)$. It follows that multiplication by x (i.e., $\{00000010\}$ or $\{02\}$) can be implemented at the byte level as a left shift and a subsequent conditional bitwise XOR with $\{1b\}$. This operation on bytes is denoted by `xtime()`. Multiplication by higher powers of x can be implemented by repeated application of `xtime()`. By adding intermediate results, multiplication by any constant can be implemented.

- To multiply by 2, we need to shift left, check the carry, if the carry is 1, the result must be XORed with `0x1b` for polynomial reduction
- To multiply by 3, after multiplying by 2, the result is XORed with the input, i.e., $03a = 02a + a$
- Many example AES codes can be found online making use of this (\rightarrow homework)

AES Key Schedule

Omitted for reasons of brevity. It is not that important for what we want to do.

AES: How to Implement

- Relatively straightforward for initial functionality
 - Performance may not be well-adapted to targeted platform
 - Very unlikely to obtain a secure implementation
- In Software:
 - Sbox can be implemented as Look Up Tables (LUT)
 - 4 Sboxes and a MixColumns merged into 8-bit to 32-bit LUT called T-table
 - Typically the fastest approach if not constrained by memory (4kb)
- In Hardware:
 - Same as software when not constrained by memory
 - Alternatively, composite field GF multipliers used for Sbox
 - In addition: various implementation styles possible for size/speed
 - Pipelined; iterated round designs: single Sbox, quarter, full

Risks of Look Up Tables ... Caches

- Most platforms are optimized for performance and low power (not security!)
- Primary two reasons why complex processors use a cache
 - Performance: accessing a data block in memory is significantly slower than a cache access
 - Power consumption: accessing a data block in memory consumes more energy than a cache access
- Caches can be implemented as
 - data cache
 - instruction cache
 - unified data and instruction cache
- Nowadays, computers typically have 3 or more levels of caching!

Timing-Dependency of Caches

- Cache hits and misses affect the timing characteristics of the implementation
- Instruction flow may be constant, but data-fetch process is not
- If the crypto algorithm is implemented with LUTs and caches exist
 - Knowing if a cache miss or hit gives away information
 - Prevention is difficult, e.g., loading LUT in cache completely
- Information leak is essential to this type of attack
- Similar to pointer leaks for software exploits, this improper caching behavior can be provoked with suitable gadgets that interfere with regular caching

Types of Cache-based Attacks

- Attacks are classified by how information leak is observed
- Trace-driven attacks
 - Trace can be recorded including hits/misses: Miss Hit Miss Miss Hit Hit ...
 - Trace can be a spy process or, e.g., measured power/emanation
- Time-driven attacks
 - Execution time is different for each run executed
 - Execution time is measured ...AND precise enough
 - Enables remote attacks over the network
- Access-driven attacks
 - The attacker tampers with the cache directly
 - By deliberately injecting suspected information or garbage into the cache, the follow-up behavior of the targeted process can be observed directly
- AES-specific attacks based on partially loaded LUTs called: “Cache-Collision”

Example: How to Run Time-Driven Attack

- Adversary model:
 - Attacker only records overall time for encryption
 - Attacker controls plaintext input
- Attack strategy:
 - Key is fixed and target of the attack
 - Collect timings $T = \{t_1, t_2, \dots, t_n\}$ for n encryptions of plaintext $P = \{p_1, p_2, \dots, p_n\}$
 - Each key byte can be attacked separately following a divide-and-conquer strategy
 - Let us consider one key byte

How to Run the Attack

- Attacker guesses a value K as key candidate for the key
 - Create two groups: group1 and group0
 - for $i = 1$ to n compute $\text{Sbox}(\text{firstbyte}(p_i) \text{ XOR } K)$ – this will be our ‘hypothesis’
 - If MSB of result is 1, the corresponding timing t_i goes to group1, otherwise to group0
 - On average, the timings of group1 should be higher than of group0
 - Averaging helps to get this information and reduce noise
 - The difference between averages of two groups is assigned guessed key byte K
 - This step needs to be repeated for all key candidates $K = 0 \dots 255$
 - The highest difference between the averages shows the most probable value candidate for the key
- You will perform this as homework assignment
- Closely following this ‘recipe’ will lead to success

Summary of AES and Timing-Dependency

- Timing-only attacks on symmetric cryptography primitives a niche topic
 - Complex devices: AES-NI or similar instruction sets offer constant runtime
 - Embedded devices: only limited number of scenarios where resistance against timing attacks is required but no resistance against power analysis is necessary
 - Aspects of timing are an integral part of power analysis
- Timing information from erroneous behavior combined with error messages can still have devastating effect, even when cryptographic primitive itself is secure!
 - “Padding Oracle” attacks, e.g., AES in Cipher-Block-Chaining (CBC) mode
 - Bleichenbacher Attack (1998), Lucky Thirteen Attack (2013), POODLE (2014)
 - Follow-up attack in 2016 enabled by initial patch in OpenSSL to fix Lucky Thirteen
- In addition: timing-dependency due to improper performance optimization has been the foundation for the famous SPECTRE microarchitectural attack (2018) which is a timing side-channel attack many thought not being practically possible

Summary of AES and Timing-Dependency

- Countermeasures:
 - Security-aware programming; including proper error-handling
 - No branch operations depending on sensitive data either directly or indirectly
 - Avoid instructions with data-dependent timing when processing secret data
 - Be aware of compiler optimizations and external libraries
- Make sure LUTs are completely cached (or not cached at all)
- Unknown cache architecture makes this more difficult (set-associative caches)
- Alternatively: avoid the LUT and compute result (better)
- Make use of dedicated hardware accelerators if available
- Limit precision of information available to attacker and access
- Generally avoid multi-tenant environments for highest security

Thank you for your attention!
Questions?